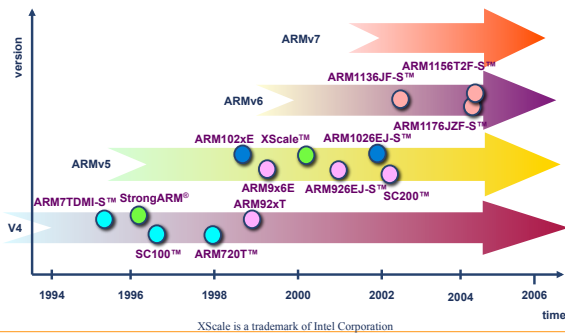


Architecture Revisions




Data Sizes and Instruction Sets

- The ARM is a 32-bit architecture.
- When used in relation to the ARM:
 - **Byte** means 8 bits
 - **Halfword** means 16 bits (two bytes)
 - **Word** means 32 bits (four bytes)
- Most ARM's implement two instruction sets
 - 32-bit ARM Instruction Set
 - 16-bit Thumb Instruction Set
- Jazelle cores can also execute Java bytecode

Processor Modes

- The ARM has seven basic operating modes:
 - **User** : unprivileged mode under which most tasks run
 - **FIQ** : entered when a high priority (fast) interrupt is raised
 - **IRQ** : entered when a low priority (normal) interrupt is raised
 - **Supervisor** : entered on reset and when a Software Interrupt instruction is executed
 - **Abort** : used to handle memory access violations
 - **Undef** : used to handle undefined instructions
 - **System** : privileged mode using the same registers as user mode

The ARM Register Set



Exception Handling

- When an exception occurs, the ARM:

- Copies CPSR into SPSR_<mode>
- Sets appropriate CPSR bits


- Change to ARM state
- Change to exception mode
- Disable interrupts (if appropriate)

- Stores the return address in LR_<mode>
- Sets PC to vector address

- To return, exception handler needs to

- Restore CPSR from SPSR_<mode>
- Restore PC from LR_<mode>

This can only be done in ARM state.



Program Status Registers



- Condition code flags**

- N = Negative result from ALU
- Z = Zero result from ALU
- C = ALU operation Carried out
- V = ALU operation Overflowed

- Sticky Overflow flag - Q flag**

- Architecture STE/J only
- Indicates if saturation has occurred

- J bit**

- Architecture STE/J only
- J = 1: Processor in Jazelle state

- Interrupt Disable bits.**

- I = 1: Disables the IRQ.
- F = 1: Disables the FIQ.

- T Bit**

- Architecture xT only
- T = 0: Processor in ARM state
- T = 1: Processor in Thumb state

- Mode bits**

- Specify the processor mode

Program Counter (r15)

- When the processor is executing in ARM state:**

- All instructions are 32 bits wide
- All instructions must be word aligned
- Therefore the pc value is stored in bits [31:2] with bits [1:0] undefined (as instruction cannot be halfword or byte aligned)

- When the processor is executing in Thumb state:**

- All instructions are 16 bits wide
- All instructions must be halfword aligned
- Therefore the pc value is stored in bits [31:1] with bit [0] undefined (as instruction cannot be byte aligned)

- When the processor is executing in Jazelle state:**

- All instructions are 8 bits wide
- Processor performs a word access to read 4 instructions at once

Conditional Execution and Flags

- ARM instructions can be made to execute conditionally by postfixing them with the appropriate condition code field.

- This improves code density and performance by reducing the number of forward branch instructions.

```

CMP    r3,#0
BEQ    skip
ADD    r0,r1,r2
skip
    
```

- By default, data processing instructions do not affect the condition code flags but the flags can be optionally set by using "S". CMP does not need "S".

```

loop
    SUBS    r1,r1,#1
    BNE    loop
    
```

decrement r1 and set flags
if Z flag clear then branch

Condition Codes

- The possible condition codes are listed below
 - Note AL is the default and does not need to be specified

Suffix	Description	Flags tested
EQ	Equal	Z=1
NE	Not equal	Z=0
CS/HS	Unsigned higher or same	C=1
CC/LO	Unsigned lower	C=0
MI	Minus	N=1
PL	Positive or Zero	N=0
VS	Overflow	V=1
VC	No overflow	V=0
HI	Unsigned higher	C=1 & Z=0
LS	Unsigned lower or same	C=0 or Z=1
GE	Greater or equal	N=V
LT	Less than	N≠V
GT	Greater than	Z=0 & N=V
LE	Less than or equal	Z=1 or N=V
AL	Always	

Conditional execution examples

C source code

```
if (x0 == 0)
{
    r1 = r1 + 1;
}
else
{
    r2 = r2 + 1;
}
```

ARM instructions

unconditional

```
CMP r0, #0
BNE else
ADD r1, r1, #1
B end
else
ADD r2, r2, #1
end
...
```

conditional

```
CMP r0, #0
ADDEQ r1, r1, #1
ADDNE r2, r2, #1
...
```

- 5 instructions
- 5 words
- 5 or 6 cycles
- 3 instructions
- 3 words
- 3 cycles

Data Processing Instructions

- Consist of :
 - Arithmetic: **ADD ADC SUB SBC RSB RSC**
 - Logical: **AND ORR EOR BIC**
 - Comparisons: **CMP CMN TST TEQ**
 - Data movement: **MOV MVN**

- These instructions only work on registers, NOT memory.

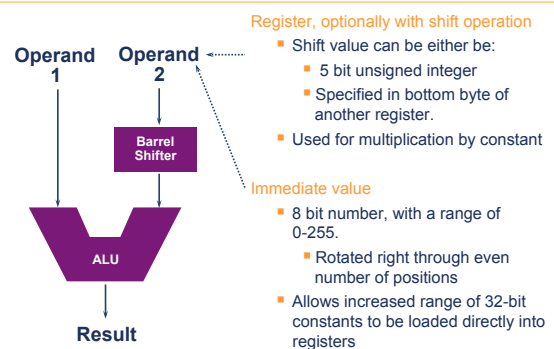
- Syntax:

<Operation>{<cond>}{S} Rd, Rn, Operand2

- Comparisons set flags only - they do not specify Rd
- Data movement does not specify Rn

- Second operand is sent to the ALU via barrel shifter.

Using a Barrel Shifter: The 2nd Operand



Data Processing Exercise


1. How would you load the two's complement representation of -1 into Register 3 using one instruction?
2. Implement an ABS (absolute value) function for a registered value using only two instructions.
3. Multiply a number by 35, guaranteeing that it executes in 2 core clock cycles.

Data Processing Solutions

1. `MOVN r6, #0`
2. `MOVS r7,r7 ; set the flags`
`RSBMI r7,r7,#0 ; if neg, r7=0-r7`
3. `ADD r9,r8,r8,LSL #2 ; r9=r8*5`
`RSB r10,r9,r9,LSL #3 ; r10=r9*7`

Immediate constants

- No ARM instruction can contain a 32 bit immediate constant
 - All ARM instructions are fixed as 32 bits long
- The data processing instruction format has 12 bits available for operand2



Quick Quiz:
`0xe3a004ff`
`MOV r0, #???`

- 4 bit rotate value (0-15) is multiplied by two to give range 0-30 in steps of 2
- Rule to remember is
 "8-bits rotated right by an even number of bit positions"

Loading 32 bit constants

- To allow larger constants to be loaded, the assembler offers a pseudo-instruction:
 - `LDR rd, =const`
- This will either:
 - Produce a `MOV` or `MOVN` instruction to generate the value (if possible).
- or
 - Generate a `LDR` instruction with a PC-relative address to read the constant from a *literal pool* (Constant data area embedded in the code).
- For example
 - `LDR r0,=0xFF` => `MOV r0,#0xFF`
 - `LDR r0,=0x55555555` => `LDR r0,[PC,#Imm12]`
 - ...
 - `DCD 0x55555555`
- This is the recommended way of loading constants into a register

Single register data transfer

LDR **STR** Word
LDRB **STRB** Byte
LDRH **STRH** Halfword
LDRSB Signed byte load
LDRSH Signed halfword load

- Memory system must support all access sizes

Syntax:

- LDR**{<cond>}{<size>} Rd, <address>
- STR**{<cond>}{<size>} Rd, <address>

e.g. **LDREQB**

Address accessed

- Address accessed by LDR/STR is specified by a base register with an offset
- For word and unsigned byte accesses, offset can be:
 - An unsigned 12-bit immediate value (i.e. 0 - 4095 bytes)
LDR r0, [r1, #8]
 - A register, optionally shifted by an immediate value
LDR r0, [r1, r2]
LDR r0, [r1, r2, LSL#2]
- This can be either added or subtracted from the base register:
LDR r0, [r1, #-8]
LDR r0, [r1, -r2, LSL#2]
- For halfword and signed halfword / byte, offset can be:
 - An unsigned 8 bit immediate value (i.e. 0 - 255 bytes)
 - A register (unshifted)
- Choice of *pre-indexed* or *post-indexed* addressing
- Choice of whether to update the base pointer (pre-indexed only)
LDR r0, [r1, #-8]!

Load/Store Exercise

Assume an array of 25 words. A compiler associates y with r1. Assume that the base address for the array is located in r2. Translate this C statement/assignment using just three instructions:

```
array[10] = array[5] + y;
```

Load/Store Exercise Solution

```
array[10] = array[5] + y;
```

```
LDR    r3, [r2, #5]    ; r3 = array[5]
ADD     r3, r3, r1      ; r3 = array[5] + y
STR     r3, [r2, #10]   ; array[5] + y = array[10]
```


Load and Store Multiples

- Syntax:
 - <LDM|STM>{<cond>}<addressing_mode> Rb{!}, <register list>

- 4 addressing modes:

- LDMIA / STMIA increment after
- LDMIB / STMIB increment before
- LDMDA / STMDA decrement after
- LDMDB / STMDB decrement before

LDMxx r10, {r0,r1,r4}
STMxx r10, {r0,r1,r4}



Multiply and Divide

- There are 2 classes of multiply - producing 32-bit and 64-bit results
- 32-bit versions on an ARM7TDMI will execute in 2 - 5 cycles

- MUL r0, r1, r2 ; r0 = r1 * r2
- MLA r0, r1, r2, r3 ; r0 = (r1 * r2) + r3


- 64-bit multiply instructions offer both signed and unsigned versions
- For these instructions there are 2 destination registers

- [U|S]MULL r4, r5, r2, r3 ; r5:r4 = r2 * r3
- [U|S]MLAL r4, r5, r2, r3 ; r5:r4 = (r2 * r3) + r5:r4

- Most ARM cores do not offer integer divide instructions
- Division operations will be performed by C library routines or inline shifts

Branch instructions

- Branch : B{<cond>} label
- Branch with Link : BL{<cond>} subroutine_label



- The processor core shifts the offset field left by 2 positions, sign-extends it and adds it to the PC
 - ± 32 Mbyte range
 - How to perform longer branches?

Register Usage


Arguments into function Result(s) from function otherwise corruptible (Additional parameters passed on stack)	Register	The compiler has a set of rules known as a Procedure Call Standard that determine how to pass parameters to a function (see AAPCS)
	r0 r1 r2 r3	
Register variables Must be preserved	r4 r5 r6 r7 r8 r9/sb r10/sl r11	CPSR flags may be corrupted by function call. Assembler code which links with compiled code must follow the AAPCS at external interfaces
	r12	
Scratch register (corruptible)	r12	The AAPCS is part of the new ABI for the ARM Architecture
Stack Pointer Link Register Program Counter	r13/sp r14/lr r15/pc	

- Stack base
- Stack limit if software stack checking selected

- SP should always be 8-byte (2 word) aligned
- R14 can be used as a temporary once value stacked

ARM Branches and Subroutines

- **B <label>**
 - PC relative. ± 32 Mbyte range.
- **BL <subroutine>**
 - Stores return address in LR
 - Returning implemented by restoring the PC from LR
 - For non-leaf functions, LR will have to be stacked



PSR access



- MRS and MSR allow contents of CPSR / SPSR to be transferred to / from a general purpose register or take an immediate value
 - MSR allows the whole status register, or just parts of it to be updated
- Interrupts can be enable/disabled and modes changed, by writing to the CPSR
 - Typically a read/modify/write strategy should be used:

```

MRS r0,CPSR      ; read CPSR into r0
BIC r0,r0,#0x80   ; clear bit 7 to enable IRQ
MSR CPSR_c,r0     ; write modified value to 'c' byte only
  
```

- In User Mode, all bits can be read but only the condition flags (`_f`) can be modified


Agenda

Introduction to ARM Ltd
Fundamentals, Programmer's Model, and Instructions

- Core Family Pipelines
- AMBA

Pipeline changes for ARM9TDMI

ARM7TDMI




ARM9TDMI




ARM10 vs. ARM11 Pipelines

ARM10



ARM11




Agenda

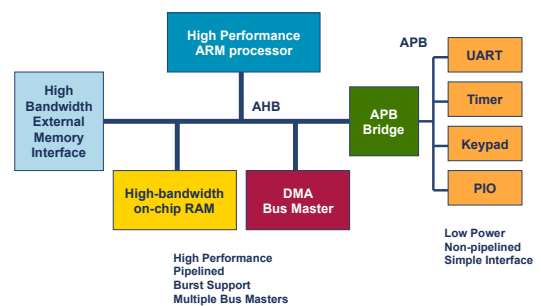
Introduction to ARM Ltd
Fundamentals, Programmer's Model, and Instructions
Core Family Pipelines

■ AMBA

Example ARM-based System



An Example AMBA System



AHB Structure

